

U SIDHARTH

GenAI Engineer | Agentic AI | LLM Systems and Automation | Business and Data Analytics

sidharthunnikrishnan2002@gmail.com | +91-9207446545 | linkedin.com/in/sidharthunnikrishnan13 | Bengaluru, Karnataka

PROFESSIONAL SUMMARY

GenAI and Agentic AI engineer with active internship experience shipping production LLM systems, self-improving AI pipelines, and multi-agent orchestration frameworks. Recently built an end-to-end automated B2B lead generation pipeline at a fintech startup using Claude AI, real-time web grounding, and hallucination safeguards. Proficient in prompt engineering, RAG pipelines, embeddings, agent loops, and tool integration using LangChain and LangGraph. Hands-on GCP experience with Pub/Sub, Dataflow, BigQuery, and Cloud Run. Strong Python and API development fundamentals with a focus on building observable, traceable, and production-safe AI systems.

TECHNICAL SKILLS

GenAI and LLMs: Prompt Engineering, Embeddings, RAG Pipelines, LLM Fine-tuning, Hallucination Safeguards, Structured Outputs, Grounding

Agentic AI: LangGraph, LangChain, Agent Loops, Tool Integration, Memory, Multi-Agent Orchestration, Human-in-the-Loop, Self-improving Pipelines

LLM APIs and Models: Anthropic Claude API (Sonnet and Haiku), OpenAI API, Groq LLaMA 3.3, Hugging Face, Embedding Models

Data and Automation: Python, Apify, GitHub Actions, Google Sheets API, SignalHire API, REST APIs, SDKs, FastAPI, Docker

GCP Services: Vertex AI, BigQuery, Cloud Storage, Pub/Sub, Cloud Run, Dataflow, GCP IAM

Databases: MongoDB Atlas, PostgreSQL, SQLite, FAISS, Pinecone, Vector Databases

ML Frameworks: Scikit-learn, TensorFlow, XGBoost, LightGBM, Regression, Classification, Clustering

Programming: Python, SQL, TypeScript, Pandas, NumPy, Git, GitHub Actions, Power BI

PROFESSIONAL EXPERIENCE

AI Engineer Intern | *Karbon Business, Bengaluru* | May 2026 – June 2026

- Built a nightly B2B lead generation pipeline using Claude AI — scraping LinkedIn profiles via Apify, classifying leads, scoring against an ICP, and enriching contact data via SignalHire API — generating 200+ qualified B2B leads
- Engineered a self-improving scoring engine that auto-updates reference profiles and recalibrates signal weights from sales team feedback in Google Sheets, improving lead quality with zero manual tuning
- Grounded all LLM scoring in real-time web data — verifying restaurant presence on DoorDash, UberEats, and Grubhub before scoring, eliminating reliance on model memory
- Applied hallucination safeguards across all Claude API calls: temperature=0, URL cross-validation, output range clamping, and grounding constraints ensuring traceable, data-bound AI decisions

Data Engineer Intern | *Dataneurus, Bengaluru* | January 2026 – March 2026

- Built GenAI-ready data pipelines on GCP using Pub/Sub, Dataflow, and BigQuery, developing hands-on experience with the full GCP services stack for real-time data ingestion and processing
- Developed a configuration-driven stream processing accelerator enabling dynamic transformations without code changes, demonstrating clean modular Python design with reusable components
- Designed and implemented a Dead Letter Queue (DLQ) replay system with retry controls, error categorisation, and status tracking — applying observability principles: logs, error handling, and retries

AI/ML Intern | *Systems Domain, Bengaluru* | June 2025 – November 2025

- Built an AI-powered Telegram tutoring chatbot using RAG, Pinecone vector search, and LLMs to help students understand educational content from pre-uploaded study materials through scalable, context-aware conversations
- Designed Power BI dashboards for KPI tracking and metric extraction, collaborating across cross-functional teams from design to production deployment

PROJECTS

Automated B2B Lead Generation Pipeline

Python, Claude AI (Sonnet and Haiku), MongoDB Atlas, Apify, SignalHire API, Google Sheets API, GitHub Actions

- Designed a multi-step nightly pipeline that scrapes LinkedIn Sales Navigator profiles via Apify, classifies ethnicity and origin using Claude AI, scores leads against a defined Ideal Customer Profile (ICP), and enriches contact data via SignalHire API
- Built a self-improving scoring engine — sales team feedback in Google Sheets automatically triggers reference profile updates and signal weight recalibration via Claude AI, progressively improving lead quality with zero manual tuning
- Grounded all AI scoring decisions in real-time web data by verifying restaurant presence on third-party delivery platforms before assigning scores, eliminating reliance on LLM memory
- Engineered production-safe LLM usage with temperature=0, URL cross-validation, output range clamping, and explicit grounding constraints across all Claude API calls

Job Hunt Agent – Autonomous Agentic AI Pipeline

Python, LangGraph, Groq LLaMA 3.3, FastAPI, Telegram, GitHub Actions, SQLite

- Architected an 8-node LangGraph agentic pipeline implementing agent loops, tool integration, and stateful memory — covering NLP query parsing, semantic scoring, resume tailoring, and output generation
- Built an LLM-powered evaluation engine using prompt engineering and embedding-based scoring to rate and filter job postings on a 0 to 100 scale, with human-in-the-loop Telegram approval checkpoints
- Automated daily execution via GitHub Actions with SQLite-backed LangGraph checkpointing for stateful, resumable runs with structured error handling and observability

NutriBot – RAG and GenAI Nutrition Assistant

Python, LangGraph, LangChain, FastAPI, FAISS, PostgreSQL, Next.js

- Implemented GenAI workflows using prompt engineering, embeddings, and RAG pipelines — a LangGraph multi-node pipeline with retrieval, context synthesis, generation, and validation nodes with automatic retry
- Built metadata-aware FAISS semantic retrieval over medical PDFs with condition-based filtering, and an automated output validation layer cross-checking every LLM generation against a structured knowledge base

SplitMint – Natural Language Parsing and GenAI Finance App

Next.js 16, TypeScript, MongoDB, Groq LLaMA 3.3, LangChain, Vercel | [splitmint.vercel.app](#)

- Developed an AI-assisted expense splitting and management app that tracks user spending through MintSense, an embedded AI assistant for logging, categorising, and querying expenses in natural language
- Built MintSense AI — a natural language query parser applying prompt engineering and semantic modelling to extract structured fields from free-text, with a token-efficient context synthesis pipeline achieving approximately 35 percent fewer tokens than JSON
- Deployed a streaming AI chatbot with time-aware memory context showcasing agent memory fundamentals and multi-turn LLM interactions in a live production application

EDUCATION

Master of Science in Data Science | Christ (Deemed to be University), Bengaluru, Karnataka | Completed

Bachelor of Science in Computer Science, Mathematics and Statistics | Kristu Jayanti Autonomous College, Bengaluru, Karnataka | Completed

CERTIFICATIONS

- AWS Cloud Foundations – AWS Academy
- Infosys Applied Generative AI Certification
- Python for Data Analytics – Pandas, NumPy, Matplotlib

LEADERSHIP AND EXTRACURRICULAR

- Co-founded the Departmental Book Club, fostering collaborative learning and peer knowledge sharing
- Active member of Tech Revamp Club — contributed to technical event management and workshop delivery
- Mentored peers and contributed to technical documentation, demonstrating communication and collaboration